

# Lesson 10 Clustering

# Outline

- ◆ Cluster Analysis and techniques
- ◆ Partitioning Algorithms: K-Means
- ◆ Determine the number of clusters
- ◆ Hierarchical Algorithms: Divisive and Agglomerative

# Cluster Analysis

- ◆ Clustering is a technique used for automatic identification of natural groupings of things
  - data instances that are similar to each other are categorized into one cluster
  - data instances that are very different from each other into different clusters.
  - Learns the clusters of things from past data, then assigns new instances to their cluster homes
- ◆ Unsupervised learning
  - There is no class label
- ◆ Also known as data segmentation


# Clustering Applications

- ◆ Group people of similar sizes together to make small, medium and large T-Shirt sizes
  - To provide fit clothes at mass production rates
- ◆ Segment customers according to their similarities
  - To do targeted marketing.
- ◆ Image Segmentation
  - Goal: Break up the image into meaningful or perceptually similar regions



# An illustration

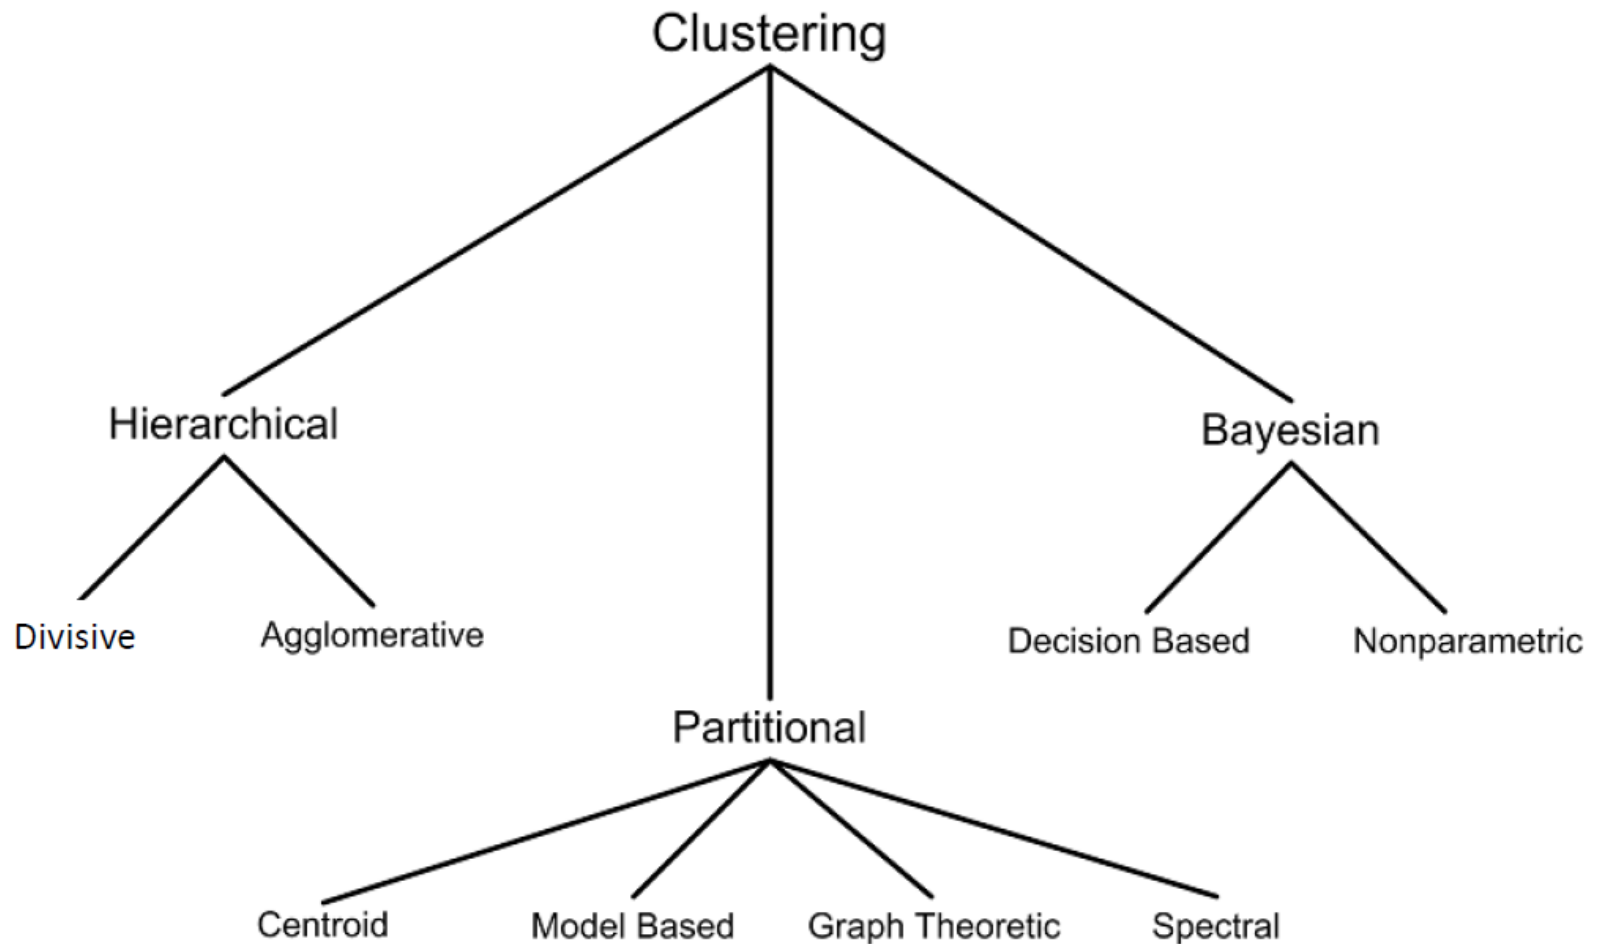
- ◆ This data set has three natural groups of data points, i.e., 3 natural clusters.



# Cluster Analysis for Data Mining

- ◆ A measure of similarity: most cluster analysis methods use a distance measure to calculate the similarity between pairs of items
  - Euclidean distance, Manhattan distance etc.
- ◆ A good clustering method will produce high quality clusters
  - high intra-class similarity: **cohesive** within clusters
  - low inter-class similarity: **distinctive** between clusters

# Clustering techniques

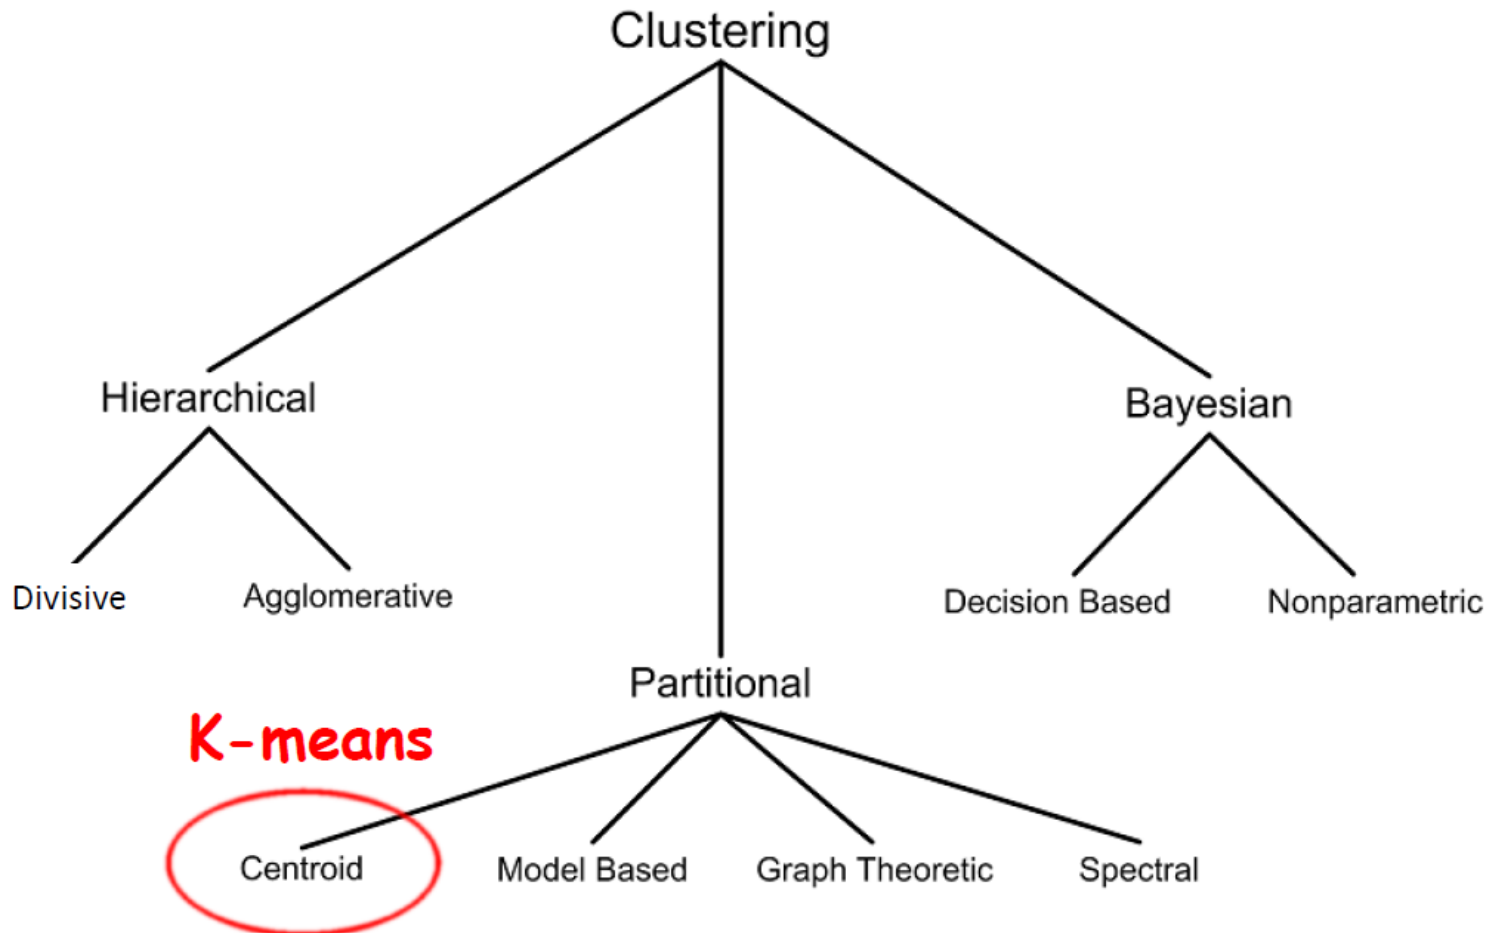


# Clustering techniques

- ◆ Hierarchical algorithms find successive clusters using previously established clusters. These algorithms can be either agglomerative (“bottom-up”) or divisive (“top-down”):
  - Agglomerative algorithms begin with each element as a separate cluster and merge them into successively larger clusters;
  - Divisive algorithms begin with the whole set and proceed to divide it into successively smaller clusters.
- ◆ Partitional algorithms: Given a set of  $n$  objects, a partitioning method constructs  $k$  partitions of the data, where each partition represents a cluster and  $k \leq n$ .



# Clustering techniques



# Partitioning Algorithms: Basic Concept

- ◆ Partitioning method: Partitions a database ***D*** of ***n*** objects into a set of ***k*** clusters, such that the sum of squared distances is minimized (where *p* is a data point and *c<sub>i</sub>* is the centroid of cluster *C<sub>i</sub>*).

$$E = \sum_{i=1}^k \sum_{p \in C_i} (p - c_i)^2$$


- ◆ This is also called *SSE (sum of squared errors)*.

# Partitioning Algorithms: Basic Concept

## ◆ More about SSE:

$$E = \sum_{i=1}^k \sum_{p \in C_i} \underbrace{(p - c_i)^2}_{\text{squared distance between an object and the centroid of its cluster}}$$

sum for all clusters



squared distance  
between an object  
and the centroid of its  
cluster

sum for all objects in  
a cluster

# k-Means Clustering Algorithm: Pseudo-code

**Algorithm** K-Means( $k, D$ )

**Input** database  $D$  of  $n$  objects,  $k$  (# of clusters)

**Output**  $n$  objects and their cluster assignments

Choose  $k$  data points as the initial centroids (cluster centers)

**repeat**

**for** each data point  $x$  in  $D$  **do**

        compute the distance from  $x$  to each centroid

        assign  $x$  to the closest centroid // a centroid represents a cluster

    re-compute the centroids using the current cluster memberships

**until** the stopping criterion is met

**return** all objects and their cluster assignments

# Stopping Criteria

- ◆ The membership assignment does not change. (We will use this criterion for demo.)
- ◆ After each reassignment,  $E$  is computed and if  $E$  falls below a predefined threshold.
- ◆ If the decrease in  $E$ , between two consecutive iterations, falls below a predefined threshold.
- ◆ Run for a predetermined number of iterations (e.g., run 20 iterations).

# Distance (Dissimilarity) Measures

## ◆ Manhattan (city block) distance

- consider two data items, A and B, each with two attribute values,  $(x1, y1)$  and  $(x2, y2)$ , respectively.

$$d(A, B) = |x1 - x2| + |y1 - y2|$$

## ◆ Euclidean distance

- consider two data items, A and B, each with two attribute values,  $(x1, y1)$  and  $(x2, y2)$ , respectively.

$$d(A, B) = \sqrt{|x1 - x2|^2 + |y1 - y2|^2}$$


## ◆ They are special cases of Minkowski distance

- In general, the Minkowski distance between the two p-dimensional data items is given by the following formula, where  $q$  is a positive integer. ( $x1, y1, z1 \dots$  are the attribute values.)



$$d(A, B) = \sqrt[q]{|x1 - x2|^q + |y1 - y2|^q + |z1 - z2|^q + \dots}$$

# Worked Example

Initial dataset,  $D = \{(1,3), (3,5), (4,1), (5, 3), (5, 4), (7, 2), (7, 4)\}$






Initial dataset



Two objects (3,5),(7,2)  
are randomly chosen as  
initial centroids

# Worked Example






Objects are assigned to the cluster with the closest centroid using Euclidean distance. For example:  
 $D((1,3), C1) = 2.83 < D((1,3), C2) = 6.08$   
so (1,3) will be in cluster with C1.

New centroids are computed.  
 $C1.x = (1+3+5)/3 = 3$   
 $C1.y = (3+4+5)/3 = 4$   
 $C2.x = (4+5+7+7)/4 = 5.75$   
 $C2.y = (1+2+3+4)/4 = 2.5$




# Worked Example



Objects are reassigned based on the distances to new centroids.  
Note that object (5,4) moved to cluster of C2.  
 $D((5,4), C1) = 2 < D((5,4), C2) = 1.68$

New centroids are computed.  
 $C1.x = (1+3)/2 = 2$   
 $C1.y = (3+5)/2 = 4$   
 $C2.x = (4+5+5+7+7)/5 = 5.6$   
 $C2.y = (1+2+3+4+4)/5 = 2.8$

# Worked Example



Objects are reassigned  
based on the distances to  
new centroids.

There is no membership  
change.

So, stop here.

# R code

#this command initializes R's random number generator to a specific  
#sequence, so set.seed to ensure reproducibility.

```
set.seed(20)
```

```
irisCluster <- kmeans(iris[, 3:4], 3)
```

```
irisCluster
```


#plot the data to see the clusters

```
irisCluster$cluster <- as.factor(irisCluster$cluster)
```

```
ggplot(iris, aes(Petal.Length, Petal.Width, color = irisCluster$cluster))  
+ geom_point()
```

# Determine the number of clusters

- ◆ In practice, we don't know "right" number of clusters.
- ◆ A simple method:  $k = \sqrt{n/2}$
- ◆ Elbow method: as we increase  $k$ , find the point where the marginal benefit in regard to SSE does not increase significantly (typically a turning point in a graph). As the graph shows, reduce in SSE is significant at first, but becomes less significant after the "turning point" at around  $k = 3.5$ . In this approach, you would use an integer close to 3.5 – either 3 or 4 – as the value for the number of clusters.



# Determine the number of clusters

- ◆ Notes about determining the number of clusters in real work:
  - There are numerous statistics to measure “SSE” that can be used with the elbow method. Still, in practice, it is not always feasible to iteratively test a large number of  $k$  values because clustering large datasets can be fairly time consuming; clustering the data repeatedly is even worse. Regardless, applications requiring the exact optimal set of clusters are fairly rare. In most clustering applications., it suffices to choose a  $k$  value based on convenience rather than strict performance requirements.

# k-Means algorithm analysis

- ◆ K-Means is the most popular clustering algorithm.

- ◆ Strengths:

- Simple: easy to understand and to implement
- Efficient: K-Means is considered a linear algorithm.


- ◆ Weaknesses:

- Initial random selection of centroids affects the results
  - ◆ Run K-Means multiple times with different initial centroids
- Applicable only when the mean of objects can be defined
  - ◆ Use the k-modes method for categorical data
- Sensitive to outliers
- Not suitable to discover clusters with arbitrary shapes


# Weaknesses of K-Means:

## Clusters of arbitrary shapes

- ◆ The K-Means algorithm is not suitable for discovering clusters that are not hyper-ellipsoids (or hyper-spheres).




Two natural clusters




K-Means clusters

# Weaknesses of K-Means: Outliers

## ◆ Undesirable clusters



## ◆ Ideal clusters

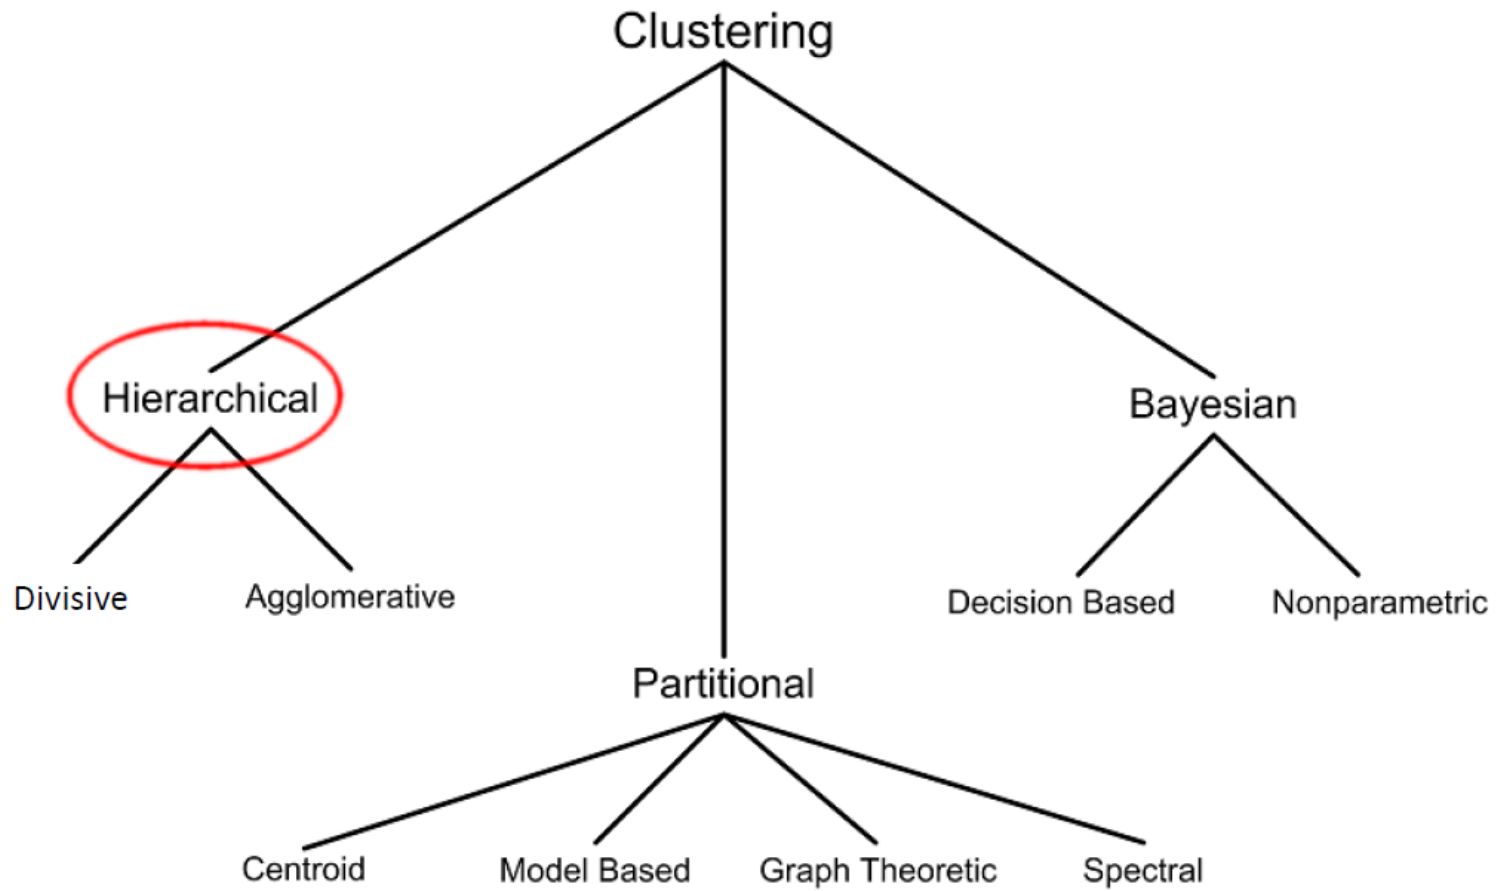




# K-Means variations

- ◆ K-Medoids – instead of mean, use medians of each cluster
  - Mean of 1, 3, 5, 7, 1009 is 205
  - Median of 1, 3, 5, 7, 1009 is 5
  - Median advantage: not affected by extreme values.

# Clustering techniques



# Hierarchical Clustering

## ◆ Agglomerative (bottom up) clustering

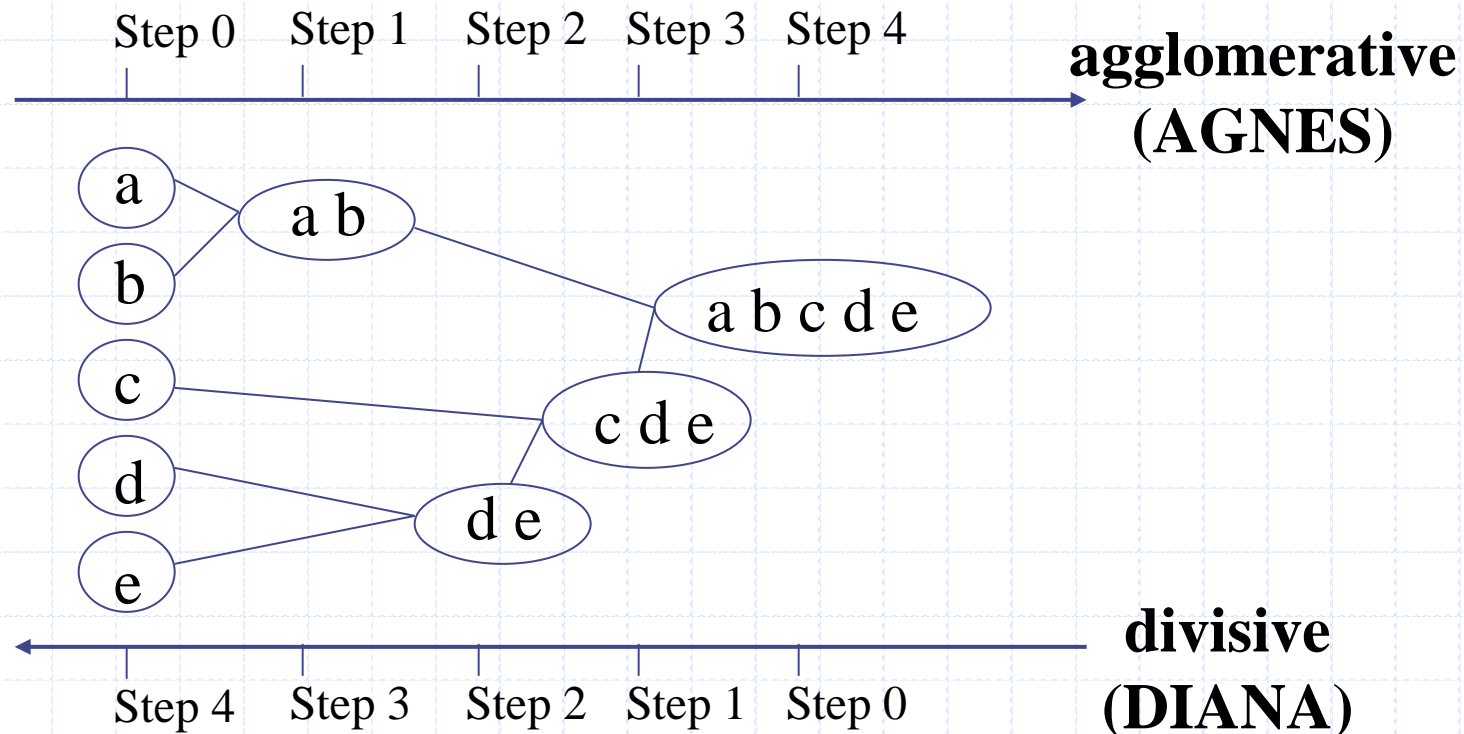
- Starting with each data point as a separate cluster
- Merging the most similar (or nearest) pair of clusters
- Stopping when all the data points are merged into a single cluster or certain termination conditions are satisfied.

## ◆ Divisive (top down) clustering

- Starting with all data points in one cluster, the root, then
- Splitting the root into a set of child cluster (according to some principle). Each child cluster is recursively divided further.
- Stopping when each cluster at the lower level is coherent enough – either containing only one object, or the objects within a cluster are sufficiently similar to each other.

- ◆ In either agglomerative or divisive hierarchical clustering, a user can specify the desired number of clusters as a termination condition<sup>27</sup>

# Hierarchical Clustering



# Distance between Clusters

## ◆ Minimum distance (single link)

- $\minDis(C_i, C_j) = \min ||x - y|| \ (x \in C_i, y \in C_j)$

## ◆ Maximum distance (complete link)

- $\maxDis(C_i, C_j) = \max ||x - y|| \ (x \in C_i, y \in C_j)$

## ◆ Average distance

- $avgDis(C_i, C_j) = \frac{1}{n_i n_j} \sum_{x \in C_i} \sum_{y \in C_j} ||x - y|| \ ((n_i, n_j \text{ are the number of data points})$

## ◆ Mean distance


- $meanDis(C_i, C_j) = ||c_i - c_j||$

$C_i, C_j$  represent clusters,  
 $c_i, c_j$  represent centroids.

# Distance between Clusters

- ◆ Minimum distance (single link): smallest distance between an element in one cluster and an element in the other.


minimum distance = 3  
(using Manhattan distance)



# Distance between Clusters

- ◆ Maximum distance (complete link): largest distance between an element in one cluster and an element in the other.

maximum distance = 10  
(using Manhattan distance)



# Distance between Clusters

◆ Average distance: average of distances between all pairs of elements.

$$d(a,e) = 6$$

$$d(a,f) = 7$$

$$d(a,g) = 8$$

$$d(b,e) = 6$$

$$d(b,f) = 5$$

$$d(b,g) = 8$$

$$d(c,e) = 7$$

$$d(c,f) = 10$$

$$d(c,g) = 7$$

$$d(d,e) = 3$$

$$d(d,f) = 6$$


$$d(d,g) = 5$$

average distance

= average of all the above

$$= 78 / 12$$

$$= 6.5$$






# Distance between Clusters

- ◆ Mean distance: distance between the centroids of two clusters

$c1 = (1.75, 4.25)$

$c2 = (6.67, 4.33)$

mean distance = 5.0  
(using Manhattan distance)



# Summary

- ◆ Despite weaknesses, K-Means is still the most popular algorithm due to its simplicity and efficiency.
- ◆ No clear evidence that any other clustering algorithm performs better in general.
- ◆ Comparing different clustering algorithms is a difficult task. No one knows the correct clusters!
- ◆ Classification and Clustering
  - Both are pattern recognition mechanisms
  - Classification is supervised learning
  - Clustering is unsupervised learning